

Deep Learning

Week 10 : Mid-Term Review + Sparse
Coding + Self-Supervised Learning

Themes

50 % : GANs + VAEs


30 % : Fully Connected + Convolutional Neural Networks

20 % : NeuroAI + Sparse Coding + Self-Supervised Learning +
Miscellaneous topics

Sample Problems (P1)

Fully Connected + CNNs

What determines the shape of the learned filter when training a neural network?




- ☐ Architecture
- ☐ Dataset at Training
- ☐ Dataset at Testing
- ☐ Regularization (eg. DropOut)
- ☐ Optimization Procedure (Learning Rule)
- ☐ Loss Function
- ☐ RAM available in CUDA Device
- ☐ Number of Epochs at Training

Sample Problems (P1)

Fully Connected + CNNs

What determines the shape of the learned filter when training a neural network?




- ☒ Architecture
- ☒ Dataset at Training
- ☐ Dataset at Testing
- ☒ Regularization (eg. DropOut)
- ☒ Optimization Procedure (Learning Rule)
- ☒ Loss Function
- ☐ RAM available in CUDA Device
- ☒ Number of Epochs at Training

Sample Problems (P2)

Fully Connected + CNNs


Write the Closed Form Expression function to compute the Average Color of an Image. What are the ideal weights of a Neural Network?




Sample Problems (P2)

Fully Connected + CNNs

Write the Closed Form Expression function to compute the Average Color of an Image. What are the ideal weights of a Neural Network?



Estimate Color
(Regression)



R: 141.30
G: 105.24
B: 64.24

Sample Problems (P3)


AutoEncoders + VAEs

Name 2 differences between the regular Auto-Encoder + Variational Auto-Encoder

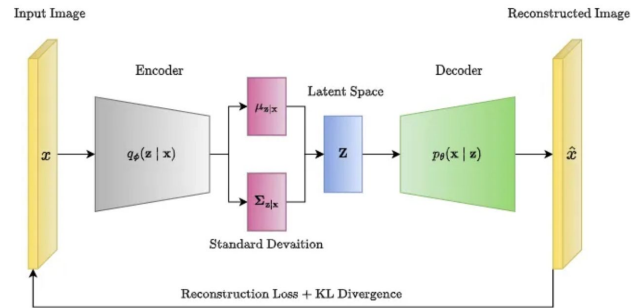
Sample Problems (P3)

AutoEncoders + VAEs

Name 2 differences between the regular Auto-Encoder + Variational Auto-Encoder



1. Deterministic vs Stochastic
2. Single-Term vs Dual-Term Loss



$$Loss_{MSE} = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{x}_i)^2$$

$$Loss_{BCE} = -\frac{1}{N} \sum_{n=1}^N \left[x_n \log \hat{x}_n + (1 - x_n) \log(1 - \hat{x}_n) \right]$$


$$\log p_{\theta}(x^{(i)}) \geq \mathcal{L}(x^{(i)}, \theta, \phi) = \underbrace{\mathbb{E}_z [\log p_{\theta}(x^{(i)} | z)]}_{\text{Reconstruct the Input Data}} - \underbrace{D_{KL}(q_{\phi}(z | x^{(i)}) || p_{\theta}(z))}_{\text{KL Divergence}}$$

$$\theta^*, \phi^* = \arg \max_{\theta, \phi} \sum_{i=1}^N \mathcal{L}(x^{(i)}, \theta, \phi)$$

Sample Problems (P3)

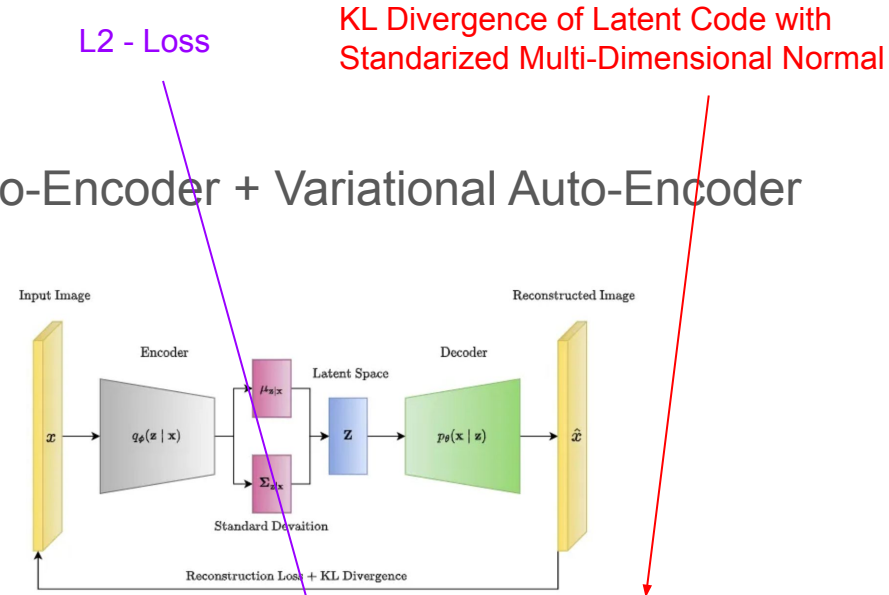
AutoEncoders + VAEs

Name 2 differences between the regular Auto-Encoder + Variational Auto-Encoder



$$Loss_{MSE} = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{x}_i)^2$$

$$Loss_{BCE} = -\frac{1}{N} \sum_{n=1}^N \left[x_n \log \hat{x}_n + (1 - x_n) \log(1 - \hat{x}_n) \right]$$



$$\log p_\theta(x^{(i)}) \geq \mathcal{L}(x^{(i)}, \theta, \phi) = \underbrace{\mathbb{E}_z [\log p_\theta(x^{(i)} | z)]}_{\text{Reconstruct the Input Data}} - \underbrace{D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z))}_{\text{KL Divergence}}$$


$$\theta^*, \phi^* = \arg \max_{\theta, \phi} \sum_{i=1}^N \mathcal{L}(x^{(i)}, \theta, \phi)$$

Recommended Reading : [Link here](#)

Sample Problems (P3)

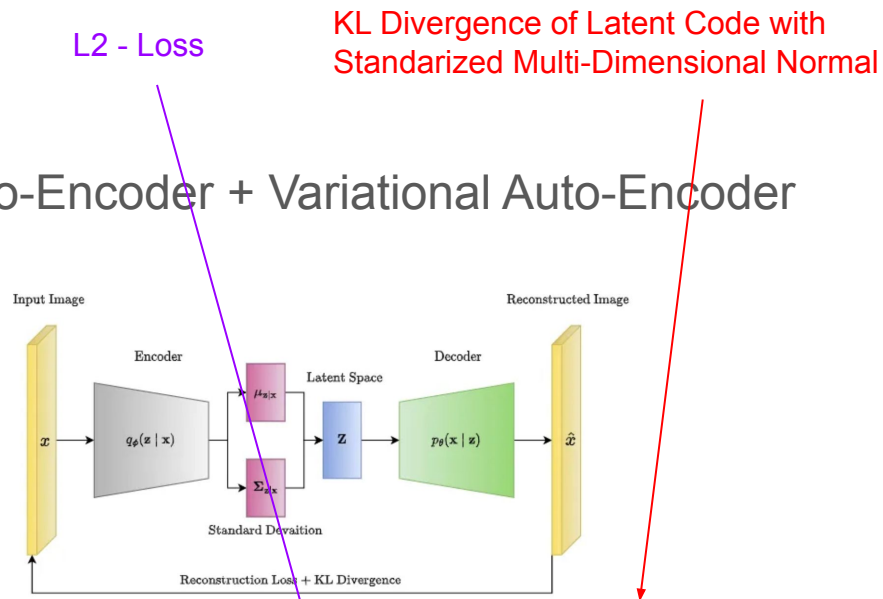
AutoEncoders + VAEs

Name 2 differences between the regular Auto-Encoder + Variational Auto-Encoder



$$Loss_{MSE} = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{x}_i)^2$$

$$Loss_{BCE} = -\frac{1}{N} \sum_{n=1}^N \left[x_n \log \hat{x}_n + (1 - x_n) \log(1 - \hat{x}_n) \right]$$



$$\log p_\theta(x^{(i)}) \geq \mathcal{L}(x^{(i)}, \theta, \phi) = \underbrace{\mathbb{E}_z [\log p_\theta(x^{(i)} | z)]}_{\text{Reconstruct the Input Data}} - \underbrace{D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z))}_{\text{KL Divergence}}$$

$$\theta^*, \phi^* = \arg \max_{\theta, \phi} \sum_{i=1}^N \mathcal{L}(x^{(i)}, \theta, \phi)$$

Sample Problems (P4)

GANs

Write 3 reasons why optimizing GAN's are difficult?

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

Sample Problems (P4)

GANs

Write 3 reasons why optimizing GAN's are difficult?

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

Min-Max Optimization Problem, where the Generator and the Discriminator are competing. Guarantees of convergence are tricky.

Highly susceptible to initialization of Generator and Discriminator

Playing with the Architecture (of G & D), Batch Size and other hyper-parameters are usually necessary to train GANs

Sample Problems (P5)

GANs + cGANs





In a Conditional GAN, the “Condition” is put in the :

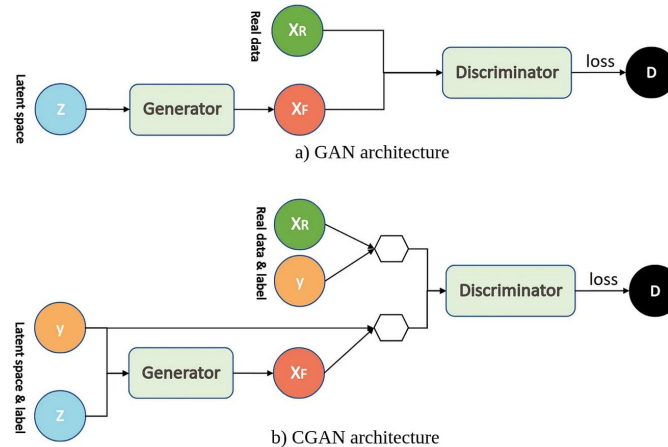
- ☐ Input of Discriminator
- ☐ Output of Discriminator
- ☐ Input of Generator
- ☐ Output of Generator

Sample Problems (P5)

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x}|\mathbf{y})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z}|\mathbf{y})))]$$

In a Conditional GAN, the “Condition” is put in the :

-  Input of Discriminator
-  Output of Discriminator
-  Input of Generator
-  Output of Generator



Sample Problems (P6)


NeuroAI

How do you compute the Gramian Matrix in a Representational Similarity Analysis (RSA) ?

Sample Problems (P6)

NeuroAI

How do you compute the Gramian Matrix in a Representational Similarity Analysis (RSA) ?



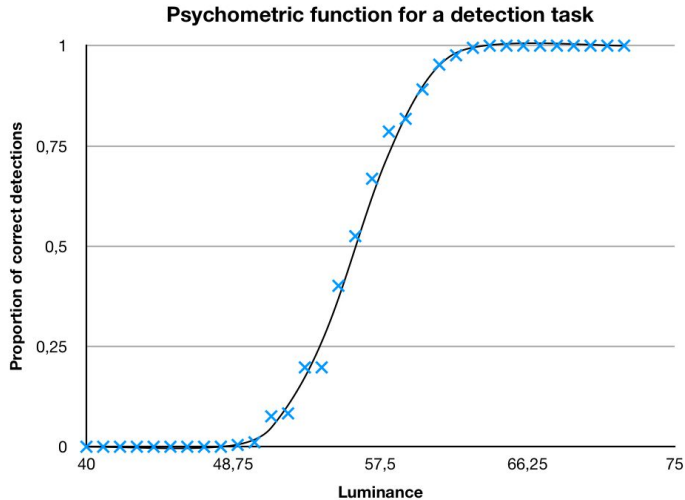


Sample Problems (P7)

Draw an Example of a Psychometric Function

Sample Problems (P7)

Draw an Example of a Psychometric Function



Sample Problems (P8)

Explain the Pipeline of the YOLO Image Classification Model

Sparse Coding Fundamentals

$$x = \text{Weights} * \text{factors} + \text{bias} + \text{noise}$$


x — vector of size u

Weights — matrix of size (k, u) , estimated model parameter

factors — vector of size k , random variables

bias — vector of size u , estimated model parameter, usually standard Gaussian

noise — vector of size u , random variables, zero-centered



Sparse Coding Fundamentals




FIG. 3 Test cases. Representative training images are shown at the left and the resulting basis functions that were learned from these examples are shown at the right. In *a*, images were composed of sparse pixels: each pixel was activated independently according to an exponential distribution, $P(x) = e^{-x}/Z$. In *b*, images were composed similarly to *a*, except with gratings instead of pixels (that is, 'sparse pixels' in the Fourier domain). In *c*, images were composed of sparse, non-orthogonal Gabor functions with the method described by Field¹². In all cases, the basis functions were initialized to random initial conditions. The learned basis functions successfully recover the sparse components from which the images were composed. The form of the sparseness cost function was $S(x) = -e^{-x^2}$, but other choices (see text) yield the same results.

Sparse codes

PCA components

Sparse Coding Fundamentals

D. Drix, V.V. Hafner and M. Schmuken / Neural Networks 131 (2020) 37–49

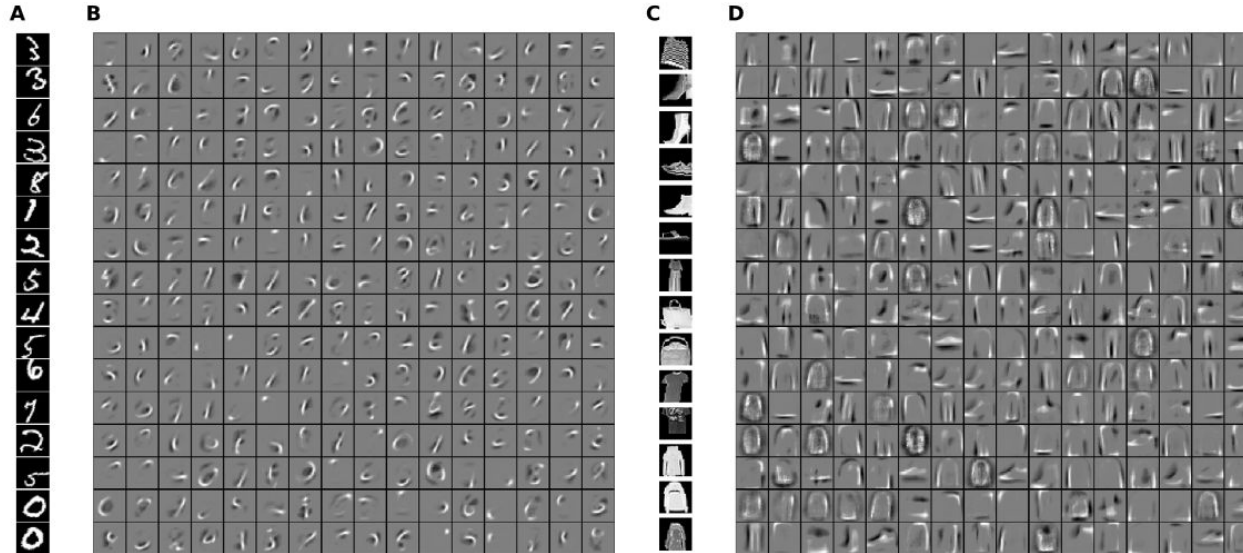
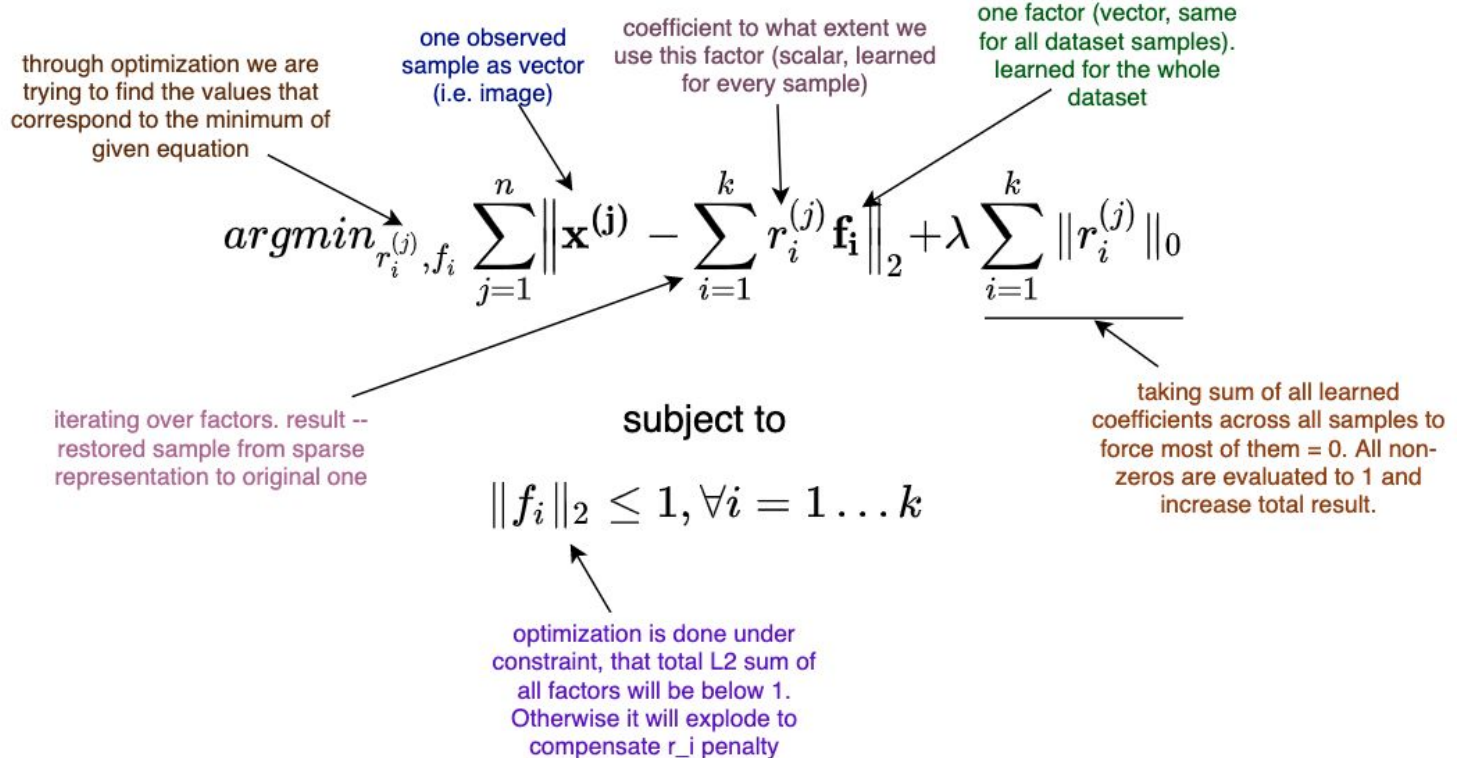


Fig. 3. The network learns independent components from the MNIST datasets. A, B: the network learns pen-stroke shapes from the MNIST dataset. A: sample input stimuli. Black corresponds to zero and white to one. B: receptive fields (weights) of a network with 256 neurons after training on 120,000 digits (28×28 pixels) with random distortions. Middle grey corresponds to zero, lighter pixels to excitatory weights, and darker pixels to inhibitory weights. C, D: the network learns the outlines and parts of the various items of clothing in the Fashion-MNIST dataset; for instance the neuron in the top-right corner of D responds to short sleeves. All other details are the same as for A and B.

Sparse Coding Fundamentals



SPARSE CODING

Topics: sparse coding

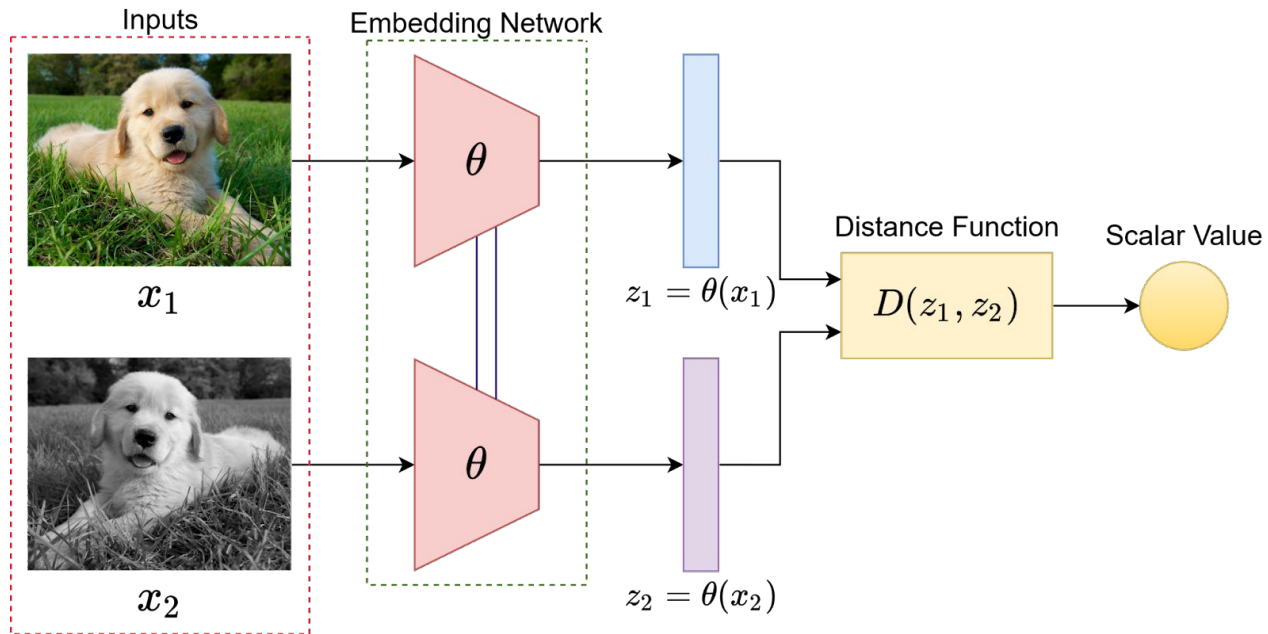
- For each $\mathbf{x}^{(t)}$ find a latent representation $\mathbf{h}^{(t)}$ such that:
 - it is sparse: the vector $\mathbf{h}^{(t)}$ has many zeros
 - we can reconstruct the original input $\mathbf{x}^{(t)}$ as well as possible
- More formally:

$$\min_{\mathbf{D}} \frac{1}{T} \sum_{t=1}^T \min_{\mathbf{h}^{(t)}} \frac{1}{2} \underbrace{\|\mathbf{x}^{(t)} - \mathbf{D} \mathbf{h}^{(t)}\|_2^2}_{\text{reconstruction error}} + \underbrace{\lambda \|\mathbf{h}^{(t)}\|_1}_{\text{sparsity penalty}}$$

reconstruction $\hat{\mathbf{x}}^{(t)}$ reconstruction vs. sparsity control

- we also constrain the columns of \mathbf{D} to be of norm 1
 - otherwise, \mathbf{D} could grow big while $\mathbf{h}^{(t)}$ becomes small to satisfy the prior
- sometimes the columns are constrained to be no greater than 1

Self-Supervised Learning



SimCLR

Barlow Twins

MoCo

A Simple Framework for Contrastive Learning of Visual Representations

Ting Chen¹ Simon Kornblith¹ Mohammad Noori¹ Geoffrey Hinton¹

Abstract

This paper presents SimCLR, a simple framework for contrastive learning of visual representations. We simply recently proposed contrastive self-supervised learning algorithms without requiring specialized architectures or a memory bank. In order to understand what enables the contrastive prediction tasks to learn useful representations, we systematically study the major components of our framework. We show that (1) the composition of data augmentations plays a critical role in defining effective predictive tasks, (2) introducing a learnable nonlinear transformation between the representation and the contrastive loss substantially improves the quality of the learned representations, and (3) contrastive learning benefits from larger batch sizes and more training steps compared to supervised learning. By combining these findings, we are able to considerably outperform previous methods for self-supervised and semi-supervised learning on ImageNet. A linear classifier trained on self-supervised representations learned by SimCLR achieves 76.5% top-1 accuracy, which is a 7% relative improvement over previous state-of-the-art, matching the performance of a supervised ResNet-50. When fine-tuned on only 1% of the labels, we achieve 85.8% top-5 accuracy, outperforming AlexNet with 100x fewer labels.¹

1. Introduction

Learning effective visual representations without human supervision is a long-standing problem. Most mainstream approaches fall into one of two classes: generative or discriminative. Generative approaches learn to generate or otherwise model pixels in the input space (Hinton et al., 2006; Kingma & Welling, 2013; Goodfellow et al., 2014).

¹Google Research, Brain Team. Correspondence to: Ting Chen <tingchen@google.com>.

Proceedings of the 37th International Conference on Machine Learning, Vienna, Austria, PMLR 119, 2020. Copyright 2020 by the authors.
Code available at <https://github.com/google-research/simclr>.




Figure 1. ImageNet Top-1 accuracy of linear classifiers trained on representations learned with different self-supervised methods (pretrained on ImageNet). Grey crosses indicate supervised ResNet-50. Our method, SimCLR, is shown in red.

However, pixel-level generation is computationally expensive and may not be necessary for representation learning. Discriminative approaches learn representations using objective functions similar to those used for supervised learning, but train networks to perform predict tasks which both the inputs and labels are derived from an unlabeled dataset. Many such approaches have relied on heuristics to design pretext tasks (Dowrich et al., 2015; Zhang et al., 2016; Norouzi & Fajos, 2016; Glotser et al., 2018), which could limit the generality of the learned representations. Discriminative approaches based on contrastive learning in the latent space have recently shown great promise, achieving state-of-the-art results (Haddad et al., 2020; Dvornitsky et al., 2014; Oord et al., 2018; Bachman et al., 2019).

In this work, we introduce a simple framework for contrastive learning of visual representations, which we call SimCLR. Not only does SimCLR outperform previous work (Figure 1), but it is also simpler, requiring neither specialized architectures (Bachman et al., 2019; Heusel et al., 2019) nor a memory bank (Wu et al., 2018; Tan et al., 2019; He et al., 2019; Miura & van der Maaten, 2019).

In order to understand what enables good contrastive representation learning, we systematically study the major components of our framework and show that:

Barlow Twins: Self-Supervised Learning via Redundancy Reduction

Jure Zbontar¹ Li Jing¹ Ishan Misra¹ Yann LeCun^{1,2} Stéphane Deny¹

Abstract

Self-supervised learning (SSL) is rapidly closing the gap with supervised methods on large computer vision benchmarks. A successful approach to SSL is to learn embeddings which are invariant to distortions of the input sample. However, a recurring issue with this approach is the existence of trivial constant solutions. Most current methods avoid such solutions by careful implementation details. We propose an objective function that naturally avoids collapse by measuring the cross-correlation matrix between the outputs of two identical networks fed with distorted versions of a sample, and making it as close to the identity matrix as possible. This causes the embedding vectors of distorted versions of a sample to be similar, while minimizing the redundancy between the components of these vectors. The method is called BARLOW TWINS, owing to neuroscience. H. Barlow's *redundancy-reduction principle* applied to a pair of identical networks. BARLOW TWINS does not require large batches nor asymmetry between the network twins such as a predictor network, gradient stopping, or a moving average on the weight updates. Intriguingly it benefits from very high-dimensional output vectors. BARLOW TWINS outperforms previous methods on ImageNet for semi-supervised classification in the low-data regime, and is on par with current state-of-the-art for ImageNet classification with a linear classifier head, and for transfer tasks of classification and object detection.¹

Figure 1. BARLOW TWINS's objective function measures the cross-correlation matrix between the embeddings of two identical networks fed with distorted versions of a batch of samples, and makes this matrix close to the identity. This causes the embedding vectors of distorted versions of a sample to be similar, while minimizing the redundancy between the components of these vectors. The method is called BARLOW TWINS, owing to neuroscience.




Figure 1. BARLOW TWINS's objective function measures the cross-correlation matrix between the embeddings of two identical networks fed with distorted versions of a batch of samples, and makes this matrix close to the identity. This causes the embedding vectors of distorted versions of a sample to be similar, while minimizing the redundancy between the components of these vectors. The method is called BARLOW TWINS, owing to neuroscience. H. Barlow's *redundancy-reduction principle* applied to a pair of identical networks. BARLOW TWINS does not require large batches nor asymmetry between the network twins such as a predictor network, gradient stopping, or a moving average on the weight updates. Intriguingly it benefits from very high-dimensional output vectors. BARLOW TWINS outperforms previous methods on ImageNet for semi-supervised classification in the low-data regime, and is on par with current state-of-the-art for ImageNet classification with a linear classifier head, and for transfer tasks of classification and object detection.¹

1. Introduction

Self-supervised learning aims to learn useful representations of the input data without relying on human annotations. Recent advances in self-supervised learning for visual data (Caron et al., 2020; Chen et al., 2020a; Grill et al., 2020; He et al., 2019; Misra & van der Maaten, 2019) show that it is a continuous, high-dimensional space and is not structured for human communication (e.g., unlike words).

Several recent studies [1, 46, 36, 35, 56, 2] present promising results on unsupervised visual representation learning using approaches related to the *contrastive loss* [29]. Though driven by various motivations, these methods can be thought of as building *dynamic dictionaries*. The “keys” (tokens) in the dictionary are sampled from data (e.g., images or patches) and are represented by an encoder network. Unsupervised learning trains encoders to perform dictionary look-up: an encoded “query” should be similar to its matching key and dissimilar to others. Learning is formulated as minimizing a contrastive loss [29].

¹Equal contribution. ²Facebook AI Research. ³New York University, NY, USA. Correspondence to: Jure Zbontar <zbontar@fb.com>, Li Jing <lingj@fb.com>, Ishan Misra <misra@fb.com>, Yann LeCun <yannl@fb.com>, Stéphane Deny <denys@openproff.com>.

Proceedings of the 36th International Conference on Machine Learning, PMLR 119, 2021. Copyright 2021 by the authors.
Code and pre-trained models on PyTorch are available at <https://github.com/facebookresearch/barlowtwins>.

Momentum Contrast for Unsupervised Visual Representation Learning

Kaiming He¹ Haoqi Fan¹ Yuxin Wu¹ Saining Xie¹ Ross Girshick¹

Facebook AI Research (FAIR)

Code: <https://github.com/facebookresearch/moco>

Abstract

We present Momentum Contrast (MoCo) for unsupervised visual representation learning. From a perspective on contrastive learning [29] at dictionary look-up, we build a *dynamic dictionary with a queue and a moving-averaged encoder*. This enables building a large and consistent dictionary *on-the-fly* that facilitates contrastive unsupervised learning. MoCo provides competitive results under the common linear protocol on ImageNet classification. More importantly, the representations learned by MoCo transfer well to downstream tasks. MoCo can outperform its supervised pre-training counterpart in 7 detection/segmentation tasks on PASCAL3D+VOC, COCO, and other datasets, sometimes surpassing it by large margins. This suggests that the gap between unsupervised and supervised representation learning has been largely closed in many visual tasks.

1. Introduction

Unsupervised representation learning is highly successful in natural language processing (e.g., as shown by GPT [50, 51] and BERT [12]), but supervised pre-training is still dominant in computer vision, where unsupervised methods generally lag behind. The reason may stem from differences in their respective signal spaces. Language tasks have discrete signal spaces (words, sub-word units, etc.) for building tokenized dictionaries, on which unsupervised learning can be based. Computer vision, in contrast, further concerns dictionary building [54, 5, 9, 3], as the raw signal is a continuous, high-dimensional space and is not structured for human communication (e.g., unlike words).

Several recent studies [1, 46, 36, 35, 56, 2] present promising results on unsupervised visual representation learning using approaches related to the *contrastive loss* [29]. Though driven by various motivations, these methods can be thought of as building *dynamic dictionaries*. The “keys” (tokens) in the dictionary are sampled from data (e.g., images or patches) and are represented by an encoder network. Unsupervised learning trains encoders to perform dictionary look-up: an encoded “query” should be similar to its matching key and dissimilar to others. Learning is formulated as minimizing a contrastive loss [29].




Figure 1. Momentum Contrast (MoCo) trains a visual representation encoder by matching an encoded query to a dictionary of encoded keys using a contrastive loss. The dictionary keys $\{k_1, k_2, \dots, k_n\}$ are defined on-the-fly by a set of data samples. The dictionary is built as a queue, with the current mini-batch encoded and the oldest mini-batch disposed, decoupling it from the mini-batch size. The keys are encoded by a slowly progressing encoder, driven by a momentum update with the query encoder. This method enables a large and consistent dictionary for learning visual representations.

From this perspective, we hypothesize that it is desirable to build dictionaries that are: (i) large and (ii) consistent as they evolve during training. Intuitively, a larger dictionary may better sample the underlying continuous, high-dimensional visual space, while the keys in the dictionary should be represented by the same or similar encoder so that their comparisons to the keys are consistent. However, existing methods that use contrastive losses can be limited in one of these two aspects (discussed later in context).

We present Momentum Contrast (MoCo) as a way of building large and consistent dictionaries for unsupervised learning with a contrastive loss (Figure 1). We maintain the dictionary as a queue of data samples; the encoded representations of the current mini-batch are enqueued, and the oldest are dequeued. The queue decouples the dictionary size from the mini-batch size, allowing it to be large. Moreover, as the dictionary keys come from the preceding several mini-batches, a *slowly progressing* key encoder, implemented as a momentum-based moving average of the query encoder, is proposed to maintain consistency.

To-Do : Leer!

A Simple Framework for Contrastive Learning of Visual Representations

Ting Chen¹, Simon Kornblith¹, Mohammad Norouzi¹, Geoffrey Hinton¹

Abstract

This paper presents SimCLR, a simple framework for contrastive learning of visual representations. We simplify recently proposed contrastive self-supervised learning algorithms without requiring specialized architectures or a memory bank. In order to understand what enables the contrastive prediction tasks to learn useful representations, we systematically study the major components of our framework. We show that (1) composition of data augmentations plays a critical role in defining effective predictive tasks, (2) introducing a learnable nonlinear transformation between the representation and the contrastive loss substantially improves the quality of the learned representations, and (3) contrastive learning benefits from larger batch sizes and more training steps compared to supervised learning. By combining these findings, we are able to considerably outperform previous methods for self-supervised and semi-supervised learning on ImageNet. A linear classifier trained on self-supervised representations learned by SimCLR achieves 76.5% top-1 accuracy, which is a 7% relative improvement over previous state-of-the-art, matching the performance of a supervised ResNet-50. When fine-tuned on only 1% of the labels, we achieve 85.8% top-5 accuracy, outperforming AlexNet with 100× fewer labels.¹

1. Introduction

Learning effective visual representations without human supervision is a long-standing problem. Most mainstream approaches fall into one of two classes: generative or discriminative. Generative approaches learn to generate or otherwise model pixels in the input space (Hinton et al., 2006; Kingma & Welling, 2013; Goodfellow et al., 2014).

¹Google Research, Brain Team. Correspondence to: Ting Chen (tingchen@google.com).

Proceedings of the 37th International Conference on Machine Learning, Vienna, Austria, PMLR 119, 2020. Copyright 2020 by the author(s).

²Code available at <https://github.com/google-research/simclr>.




Figure 1. ImageNet Top-1 accuracy of linear classifiers trained on representations learned with different self-supervised methods (trained on ImageNet). Gray cross indicates supervised ResNet-50. Our method, SimCLR, is shown in bold.

However, pixel-level generation is computationally expensive and may not be necessary for representation learning. Discriminative approaches learn representations using objective functions similar to those used for supervised learning, but train networks to perform pretext tasks where both the inputs and labels are derived from an unlabeled dataset. Many such approaches have relied on heuristics to design pretext tasks (Doersch et al., 2015; Zhang et al., 2016; Norouzi & Fares, 2016; Gidaris et al., 2018), which could limit the generality of the learned representations. Discriminative approaches based on contrastive learning in the latent space have recently shown great promise, achieving state-of-the-art results (Hadsell et al., 2006; Dosovitskiy et al., 2014; Oord et al., 2018; Bachman et al., 2019).

In this work, we introduce a simple framework for contrastive learning of visual representations, which we call SimCLR. Not only does SimCLR outperform previous work (Figure 1), but it is also simpler, requiring neither specialized architectures (Bachman et al., 2019; Hestuff et al., 2019) nor a memory bank (Wu et al., 2018; Tan et al., 2019; He et al., 2019; Misra & van der Maaten, 2019).

In order to understand what enables good contrastive representation learning, we systematically study the major components of our framework and show that:

SimCLR

How do we later do
Supervised Learning?

← → ↺ https://github.com/vturrissi/solo-learn

vturrissi / solo-learn

<> Code Issues 12 Pull requests 5 Actions Projects Security Insights

solo-learn Public

Watch 11 Fork 185 Star 1.4k

main 9 Branches 20 Tags

Go to file Add file <> Code

vturrissi Update README.md (#384) b69b4bd · 9 months ago 378 Commits

github	Fix tests on Python 3.10 (#347)	last year
docs	Add All4One (#382)	9 months ago
downstream/object_detection	Use syntax updated with Python 3.8+ (#343)	last year
scripts	Add All4One (#382)	9 months ago
solo	Add All4One (#382)	9 months ago
tests	Add All4One (#382)	9 months ago
zoo	reupload barlow checkpoint for imagenet	last year
.codecov.yml	Ci: precommit (#316)	2 years ago
.gitignore	Move from Argparse to Omegaconf (#300)	2 years ago
.pre-commit-config.yaml	[pre-commit.ci] pre-commit autoupdate (#380)	10 months ago
.readthedocs.yml	Update .readthedocs.yml (#379)	10 months ago
LICENSE	Ci: precommit (#316)	2 years ago
README.md	Update README.md (#384)	9 months ago
logo.png	Added logo and updated copyright licenses (#317)	2 years ago
main_knn.py	Fix checkpoint loading for KNN eval (#376)	last year
main_linear.py	Moved to lightning (#370)	last year

About

solo-learn: a library of self-supervised methods for visual representation learning powered by Pytorch Lightning

pytorch dino mae moco self-supervised-learning byol nvidia-dali transformer-models pytorch-lightning simclr contrastive-learning deepcluster simsiam sway barlow-twins vicreg rncir resst viboreg masked-input-prediction

Readme MIT license Code of conduct Activity 1.4k stars 11 watching 185 forks Report repository

Releases 18

1.0.8 (Latest) on Dec 7, 2023 + 17 releases

Packages

Librería de Self-Supervised Learning : Solo-Learn